

S3 - HOPE logistic regression-V2

December 4, 2020

0.1 Import data from DB.

```
[1]: import pandas as pd
import numpy as np
```

```
[2]: dfOrg = pd.read_csv('hope_dataset_cleaned.csv')

print(dfOrg.shape[0])
```

1243

```
[3]: dfOrg.head(10)
```

```
[3]:  pedido.data.attributes.age  pedido.data.attributes.diagnostic_main  \
0                75.0                FISTULA PERITONEAL
1                75.0                FISTULA PERITONEAL
2                75.0                FISTULA PERITONEAL
3                75.0                FISTULA PERITONEAL
4                75.0                FISTULA PERITONEAL
5                75.0                FISTULA PERITONEAL
6                75.0                FISTULA PERITONEAL
7                75.0                FISTULA PERITONEAL
8                75.0                FISTULA PERITONEAL
9                75.0                FISTULA PERITONEAL

    pedido.data.attributes.gender  articulo  respuesta.articlesRevisedYear  \
0                male  27395425                2018
1                male  28560554                2018
2                male  28641726                2017
3                male  26245344                2016
4                male  28942543                2018
5                male  24782153                2014
6                male  28002229                2018
7                male  27505109                2017
8                male  24850546                2015
9                male  29371050                2019

    respuesta.articlesRevisedMonth  \
```

0	1
1	4
2	12
3	12
4	6
5	6
6	9
7	4
8	1
9	4

	respuesta.pubmed_keys	utilidad
0	Abdomen,Adenocarcinoma,Antiemetics,Blood Cultu...	1.0
1	Abdomen,Adenocarcinoma,Antiemetics,Blood Cultu...	NaN
2	Abdomen,Adenocarcinoma,Antiemetics,Blood Cultu...	NaN
3	Abdomen,Adenocarcinoma,Antiemetics,Blood Cultu...	NaN
4	Abdomen,Adenocarcinoma,Antiemetics,Blood Cultu...	NaN
5	Abdomen,Adenocarcinoma,Antiemetics,Blood Cultu...	NaN
6	Abdomen,Adenocarcinoma,Antiemetics,Blood Cultu...	NaN
7	Abdomen,Adenocarcinoma,Antiemetics,Blood Cultu...	NaN
8	Abdomen,Adenocarcinoma,Antiemetics,Blood Cultu...	NaN
9	Abdomen,Adenocarcinoma,Antiemetics,Blood Cultu...	NaN

Expand pubmed_keys attribute

```
[4]: dfOrg['respuesta.pubmed_keys'] = dfOrg['respuesta.pubmed_keys'].apply(lambda x :
    → str(x).split(','))

dfOrg = dfOrg.explode('respuesta.pubmed_keys').reset_index(drop=True)

dfOrg.head(10)
```

```
[4]: pedido.data.attributes.age pedido.data.attributes.diagnostic_main \
0 75.0 FISTULA PERITONEAL
1 75.0 FISTULA PERITONEAL
2 75.0 FISTULA PERITONEAL
3 75.0 FISTULA PERITONEAL
4 75.0 FISTULA PERITONEAL
5 75.0 FISTULA PERITONEAL
6 75.0 FISTULA PERITONEAL
7 75.0 FISTULA PERITONEAL
8 75.0 FISTULA PERITONEAL
9 75.0 FISTULA PERITONEAL

pedido.data.attributes.gender articulo respuesta.articlesRevisedYear \
0 male 27395425 2018
1 male 27395425 2018
```

2	male	27395425	2018
3	male	27395425	2018
4	male	27395425	2018
5	male	27395425	2018
6	male	27395425	2018
7	male	27395425	2018
8	male	27395425	2018
9	male	27395425	2018

	respuesta.articlesRevisedMonth	respuesta.pubmed_keys	utilidad
0	1	Abdomen	1.0
1	1	Adenocarcinoma	1.0
2	1	Antiemetics	1.0
3	1	Blood Culture	1.0
4	1	Catharsis	1.0
5	1	Diuresis	1.0
6	1	Fistula	1.0
7	1	Gastrectomy	1.0
8	1	Incisional Hernia	1.0
9	1	Intestines	1.0

0.2 Transform (factorize) from Categories to continuous atributes

Transform 'pedido.data.attributes.diagnostic_main' attribute

```
[5]: dataDiagnosticMain, categoriesDiagnosticMain = pd.factorize(dfOrg['pedido.data.
    ↳attributes.diagnostic_main'])

dfOrg['pedido.data.attributes.diagnostic_main'] = dataDiagnosticMain
```

Transform 'gender' attribute

```
[6]: dataGender, categoriesGender = pd.factorize(dfOrg['pedido.data.attributes.
    ↳gender'])

dfOrg['pedido.data.attributes.gender'] = dataGender
```

Transform 'respuesta.pubmed_keys' attribute

```
[7]: categoriesORGPubMedKeys = dfOrg['respuesta.pubmed_keys'].value_counts()

print("total: " + str(categoriesORGPubMedKeys.size))
```

total: 353

```
[8]: dataPubMedKeys, categoriesPubMedKeys = pd.factorize(dfOrg['respuesta.
    ↳pubmed_keys'])
```

```
dfOrg['respuesta.pubmed_keys'] = dataPubMedKeys
```

```
[9]: dfOrg.head(10)
```

```
[9]:
```

	pedido.data.attributes.age	pedido.data.attributes.diagnostic_main	\
0	75.0	0	
1	75.0	0	
2	75.0	0	
3	75.0	0	
4	75.0	0	
5	75.0	0	
6	75.0	0	
7	75.0	0	
8	75.0	0	
9	75.0	0	

	pedido.data.attributes.gender	articulo	respuesta.articlesRevisedYear	\
0	0	27395425	2018	
1	0	27395425	2018	
2	0	27395425	2018	
3	0	27395425	2018	
4	0	27395425	2018	
5	0	27395425	2018	
6	0	27395425	2018	
7	0	27395425	2018	
8	0	27395425	2018	
9	0	27395425	2018	

	respuesta.articlesRevisedMonth	respuesta.pubmed_keys	utilidad
0	1	0	1.0
1	1	1	1.0
2	1	2	1.0
3	1	3	1.0
4	1	4	1.0
5	1	5	1.0
6	1	6	1.0
7	1	7	1.0
8	1	8	1.0
9	1	9	1.0

```
[10]: print("age NaN => " + str(dfOrg[pd.isnull(dfOrg['pedido.data.attributes.age'])].
        ↳shape[0]))
print("diagnostic_main NaN => " + str(dfOrg[pd.isnull(dfOrg['pedido.data.
        ↳attributes.diagnostic_main'])].shape[0]))
print("gender NaN => " + str(dfOrg[pd.isnull(dfOrg['pedido.data.attributes.
        ↳gender'])].shape[0]))
print("articulo NaN => " + str(dfOrg[pd.isnull(dfOrg['articulo'])].shape[0]))
```

```

print("articlesRevisedYear NaN => " + str(dfOrg[pd.isnull(dfOrg['respuesta.
↳articlesRevisedYear'])].shape[0]))
print("articlesRevisedMonth NaN => " + str(dfOrg[pd.isnull(dfOrg['respuesta.
↳articlesRevisedMonth'])].shape[0]))
print("pubmed_keys NaN => " + str(dfOrg[pd.isnull(dfOrg['respuesta.
↳pubmed_keys'])].shape[0]))
print("utilidad NaN => " + str(dfOrg[pd.isnull(dfOrg['utilidad'])].shape[0]))

```

```

age NaN => 10
diagnostic_main NaN => 0
gender NaN => 0
articulo NaN => 0
articlesRevisedYear NaN => 0
articlesRevisedMonth NaN => 0
pubmed_keys NaN => 0
utilidad NaN => 14758

```

Remove row with age eq NaN

```
[11]: dfOrg = dfOrg[pd.notnull(dfOrg['pedido.data.attributes.age'])]
```

0.3 Standardize the Data

```

[12]: from sklearn.preprocessing import StandardScaler

features = dfOrg.columns.drop(['utilidad'])

# Separating out the features
x = dfOrg.loc[:, features].values

featuresTransformed = StandardScaler().fit_transform(x)

dfStandarized = pd.DataFrame(featuresTransformed, index=dfOrg.index,
↳columns=features)
dfStandarized['utilidad'] = dfOrg['utilidad']

dfStandarized.head(10)

```

```

[12]:    pedido.data.attributes.age  pedido.data.attributes.diagnostic_main \
0                1.285887                -1.503163
1                1.285887                -1.503163
2                1.285887                -1.503163
3                1.285887                -1.503163
4                1.285887                -1.503163
5                1.285887                -1.503163
6                1.285887                -1.503163
7                1.285887                -1.503163

```

8	1.285887	-1.503163
9	1.285887	-1.503163

	pedido.data.attributes.gender	articulo	respuesta.articlesRevisedYear	\
0	0.0	-0.00421	0.633249	
1	0.0	-0.00421	0.633249	
2	0.0	-0.00421	0.633249	
3	0.0	-0.00421	0.633249	
4	0.0	-0.00421	0.633249	
5	0.0	-0.00421	0.633249	
6	0.0	-0.00421	0.633249	
7	0.0	-0.00421	0.633249	
8	0.0	-0.00421	0.633249	
9	0.0	-0.00421	0.633249	

	respuesta.articlesRevisedMonth	respuesta.pubmed_keys	utilidad
0	-1.463658	-1.089722	1.0
1	-1.463658	-1.080463	1.0
2	-1.463658	-1.071203	1.0
3	-1.463658	-1.061944	1.0
4	-1.463658	-1.052684	1.0
5	-1.463658	-1.043424	1.0
6	-1.463658	-1.034165	1.0
7	-1.463658	-1.024905	1.0
8	-1.463658	-1.015646	1.0
9	-1.463658	-1.006386	1.0

0.4 Separe data by utilidad is defined

```
[13]: dfDataSetComplete = dfStandarized[pd.notnull(dfOrg['utilidad'])]

print(dfDataSetComplete.shape[0])

dfDataSetToPredict = dfStandarized[pd.isnull(dfOrg['utilidad'])]

print(dfDataSetToPredict.shape[0])
```

830
14748

```
[14]: dfDataSetComplete.head(10)
```

	pedido.data.attributes.age	pedido.data.attributes.diagnostic_main	\
0	1.285887	-1.503163	
1	1.285887	-1.503163	
2	1.285887	-1.503163	
3	1.285887	-1.503163	

4	1.285887	-1.503163
5	1.285887	-1.503163
6	1.285887	-1.503163
7	1.285887	-1.503163
8	1.285887	-1.503163
9	1.285887	-1.503163

	pedido.data.attributes.gender	articulo	respuesta.articlesRevisedYear	\
0	0.0	-0.00421	0.633249	
1	0.0	-0.00421	0.633249	
2	0.0	-0.00421	0.633249	
3	0.0	-0.00421	0.633249	
4	0.0	-0.00421	0.633249	
5	0.0	-0.00421	0.633249	
6	0.0	-0.00421	0.633249	
7	0.0	-0.00421	0.633249	
8	0.0	-0.00421	0.633249	
9	0.0	-0.00421	0.633249	

	respuesta.articlesRevisedMonth	respuesta.pubmed_keys	utilidad
0	-1.463658	-1.089722	1.0
1	-1.463658	-1.080463	1.0
2	-1.463658	-1.071203	1.0
3	-1.463658	-1.061944	1.0
4	-1.463658	-1.052684	1.0
5	-1.463658	-1.043424	1.0
6	-1.463658	-1.034165	1.0
7	-1.463658	-1.024905	1.0
8	-1.463658	-1.015646	1.0
9	-1.463658	-1.006386	1.0

0.5 Logistic Regression

```
[15]: from sklearn import linear_model
from sklearn.model_selection import train_test_split
from sklearn import model_selection
from sklearn.metrics import accuracy_score
```

```
[16]: dfDataSetComplete.describe()
```

	pedido.data.attributes.age	pedido.data.attributes.diagnostic_main	\
count	830.000000	830.000000	
mean	0.323607	-0.039986	
std	1.006114	0.650873	
min	-1.554838	-1.503163	
25%	-1.006628	-0.586347	
50%	0.737677	0.101264	

75%	1.236049	0.215866
max	1.584910	1.820293

	pedido.data.attributes.gender	articulo \
count	830.0	830.000000
mean	0.0	0.076390
std	0.0	0.742731
min	0.0	-1.880409
25%	0.0	-0.309977
50%	0.0	0.301868
75%	0.0	0.553395
max	0.0	1.215055

	respuesta.articlesRevisedYear	respuesta.articlesRevisedMonth \
count	830.000000	830.000000
mean	0.093236	-0.168806
std	1.156247	1.082831
min	-2.656375	-1.463658
25%	-0.189157	-1.178433
50%	0.427648	-0.893208
75%	1.044452	1.103364
max	1.044452	1.673814

	respuesta.pubmed_keys	utilidad
count	830.000000	830.000000
mean	-0.031101	0.583133
std	0.925381	0.493338
min	-1.089722	0.000000
25%	-0.904532	0.000000
50%	-0.358219	1.000000
75%	0.808482	1.000000
max	2.049259	1.000000

We check the number of results

```
[17]: dfDataSetComplete.groupby('utilidad').size()
```

```
[17]: utilidad
0.0    346
1.0    484
dtype: int64
```

Chooosed “age”, “diagnostic_main”, “month” and “pubmed_keys” attributes (based on PCA_V3 study)

```
[18]: dataToTrain = dfDataSetComplete[["pedido.data.attributes.age",
    "pedido.data.attributes.diagnostic_main",
    "respuesta.articlesRevisedMonth",
```



```

        "respuesta.pubmed_keys",
        "utilidad"
    ]]

X = np.array(dataToTrain.drop(['utilidad'],1))
y = np.array(dataToTrain['utilidad'])
X.shape

```

[18]: (830, 4)

```
[19]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

```
[20]: model = linear_model.LogisticRegression()
model.fit(X_train,y_train)

model.score(X_train,y_train)
```

[20]: 0.5627009646302251

```
[21]: # see: https://es.wikipedia.org/wiki/Validaci%C3%B3n\_cruzada
kfold = model_selection.KFold(n_splits=10)
cv_results = model_selection.cross_val_score(model, X_train, y_train, cv=kfold,
    ↳scoring='accuracy')
msg = "%s: %f (%f)" % ("Logistic Regression", cv_results.mean(), cv_results.
    ↳std())
print(msg)
```

Logistic Regression: 0.554531 (0.059046)

```
[22]: predictions = model.predict(X_test)
print(accuracy_score(y_test, predictions))
```


0.49038461538461536

```
[23]: import seaborn as sn
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix
```

```
[24]: cf = confusion_matrix(y_test, predictions)

df_cm = pd.DataFrame(cf, range(2), range(2))
sn.set(font_scale=1.4) # for label size
sn.heatmap(df_cm, annot=True, annot_kws={"size": 16}) # font size

plt.show()
```



0.6 Run Prediction

```
[25]: result = model.predict(dfDataSetToPredict[["pedido.data.attributes.age",
        "pedido.data.attributes.diagnostic_main",
        "respuesta.articlesRevisedMonth",
        "respuesta.pubmed_keys"
    ]])

result
```

```
[25]: array([0., 0., 0., ..., 1., 1., 1.])
```

0.7 Try with all atributes

```
[26]: X = np.array(dfDataSetComplete.drop(['utilidad'],1))
y = np.array(dfDataSetComplete['utilidad'])
X.shape
```

```
[26]: (830, 7)
```

```
[27]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
model = linear_model.LogisticRegression()
model.fit(X_train,y_train)

model.score(X_train,y_train)
```

[27]: 0.5530546623794212

```
[28]: kfold = model_selection.KFold(n_splits=10)
cv_results = model_selection.cross_val_score(model, X_train, y_train, cv=kfold,
→scoring='accuracy')
msg = "%s: %f (%f)" % ("Logistic Regression", cv_results.mean(), cv_results.
→std())
print(msg)
```

Logistic Regression: 0.541654 (0.047477)


```
[29]: predictions = model.predict(X_test)
print(accuracy_score(y_test, predictions))
```

0.5

```
[30]: cf = confusion_matrix(y_test, predictions)

df_cm = pd.DataFrame(cf, range(2), range(2))
sn.set(font_scale=1.4) # for label size
sn.heatmap(df_cm, annot=True, annot_kws={"size": 16}) # font size

plt.show()
```



[]: